

# logistic Regression

Data set: Titanic Data Set from Kaggle. Goal: trying to predict a classification- survival or deceased.

Note: Data is a "semi-cleaned" version of the titanic data set. It has non values.

## Import Libraries

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

## Load the data set

```
In [30]: train_org = pd.read_csv('Data/titanic_train.csv')
train_org.head(5)
```

Out[30]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	C
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	

```
In [90]: # data statistic
print(train_org.shape)
print('total nun values= ',train_org.isnull().sum().sum())
print(train_org.isnull().sum())
```


```
(891, 12)
total nun values= 866
PassengerId      0
Survived         0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        2
dtype: int64
```

```
In [100]: # object columns
s = (train_org.dtypes == 'object')
#print (s)
col_object = list(s[s].index)
print(col_object)
```

```
['Name', 'Sex', 'Ticket', 'Cabin', 'Embarked']
```

```
In [34]: sns.heatmap(train_org.isnull(), cmap='viridis')
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x210860c1f98>
```



## Preprocessing on Data

In Cabin column, around 77% of data are missing --> drop the column In Age column, around 20% of the data are missing --> impute ...

Categorical column like Sex must be converted into numerical values

Column Ticket and Name will also be dropped, (no useful features)

```
In [19]: # make a copy from original data
train_clean = train_org.copy()
```

### Drop un-necessary columns

```
In [22]: print(train_clean.shape)
train_clean.columns

(891, 12)
```

```
Out[22]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
               'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
              dtype='object')
```

```
In [24]: train_clean.drop(columns=['Name', 'Ticket', 'Cabin'], inplace=True)
```

```
In [25]: print(train_clean.shape)
train_clean.columns

(891, 9)
```

```
Out[25]: Index(['PassengerId', 'Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch',
               'Fare', 'Embarked'],
              dtype='object')
```

## Handle missing Data in Age column

Filling in the mean age of all the passengers (imputation).

Or here, fill them by the average age of the passenger in class.

```
In [31]: print('Average age of all pasenger = %d' %(train_clean['Age'].dropna().mean(
)))
```

Average age of all pasenger = 29

```
In [45]: train_clean.corrwith(train_clean['Age'])
```

```
Out[45]: PassengerId    0.036847  
Survived    -0.077221  
Pclass     -0.369226  
Age         1.000000  
SibSp      -0.308247  
Parch      -0.189119  
Fare       0.096067  
dtype: float64
```


```
In [52]: train_clean.groupby('Pclass')['Age'].mean().reset_index()
```

```
Out[52]:
```

	Pclass	Age
0	1	38.233441
1	2	29.877630
2	3	25.140620

```
In [44]: sns.boxenplot(x='Pclass', y='Age', data=train_clean)
```

```
Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x2108c4f2d68>
```



```
In [53]: def impute_age(cols):
    Age = cols[0]
    Pclass = cols[1]

    if pd.isnull(Age):

        if Pclass == 1:
            return 38

        elif Pclass == 2:
            return 30

        else:
            return 25

    else:
        return Age
```

```
In [55]: train_clean['Age'] = train_clean[['Age', 'Pclass']].apply(impute_age, axis=1)
```

```
In [56]: train_clean['Age'].isnull().sum()
```

```
Out[56]: 0
```

## Handle numerical column

column Sex, Embarked

```
In [59]: sex = pd.get_dummies(train_clean['Sex'], drop_first=True)
```

```
In [63]: train_clean['Embarked'].value_counts()
```

```
Out[63]: S    644
        C    168
        Q     77
        Name: Embarked, dtype: int64
```

```
In [115]: embarked = pd.get_dummies(train_clean['Embarked'], drop_first=True)
```

```
In [116]: train_clean.drop(columns=['Sex', 'Embarked'], inplace=True, axis=1)
```

```
In [117]: print(sex.shape)
        print(embarked.shape)
        print(train_clean.shape)
```

```
(891, 1)
(891, 2)
(891, 7)
```

```
In [118]: train = pd.concat([train_clean, sex, embarked], axis=1)
```

```
In [119]: train.shape
```

```
Out[119]: (891, 10)
```

```
In [120]: train.columns
```

```
Out[120]: Index(['PassengerId', 'Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare',  
                'male', 'Q', 'S'],  
                dtype='object')
```

```
In [124]: (train.dtypes == 'Object').any()
```

```
Out[124]: False
```

## Logistic Regression Model

```
In [125]: X = train.drop(columns=['Survived'], axis=1)  
y = train['Survived']
```

```
In [126]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30)
```

```
In [127]: from sklearn.linear_model import LogisticRegression
```

```
In [132]: def get_acc (max_iter):  
    log_m = LogisticRegression(max_iter=max_iter)  
    log_m.fit(X_train, y_train)  
    log_m.predict(X_test)  
    acc_train = log_m.score(X_train, y_train)  
    acc_test = log_m.score(X_test, y_test)  
    return [acc_train, acc_test]
```

```
In [135]: import warnings; warnings.simplefilter('ignore')
```


```
In [136]: get_acc(100)
```

```
Out[136]: [0.7929373996789727, 0.835820895522388]
```

```
In [143]: train_acc = []  
test_acc = []  
iteration = [10, 50, 100, 200, 500]  
  
for i in iteration:  
    res = get_acc(i)  
    train_acc.append(res[0])  
    test_acc.append(res[1])
```

```
In [144]: plt.plot(iteration, train_acc)
plt.plot(iteration, test_acc)
plt.legend(['train_acc', 'test_acc'])
```

Out[144]: <matplotlib.legend.Legend at 0x2108faa4400>



In [ ]:

## Random Forest model

```
In [146]: from sklearn.ensemble import RandomForestClassifier
```


```
In [149]: def get_acc_rfc (min_samples_split):
    mdl = RandomForestClassifier(min_samples_split=min_samples_split)
    mdl.fit(X_train, y_train)
    #log_m.predict(X_test)
    acc_train = mdl.score(X_train, y_train)
    acc_test = mdl.score(X_test, y_test)
    return [acc_train, acc_test]
```

```
In [150]: train_acc = []
test_acc = []
min_sample = [2, 5, 10, 50, 100]

for i in min_sample:
    res = get_acc_rfc(i)
    train_acc.append(res[0])
    test_acc.append(res[1])
```

```
In [152]: plt.plot(min_sample, train_acc)
plt.plot(min_sample, test_acc)
plt.legend(['train_acc', 'test_acc'])
```

Out[152]: <matplotlib.legend.Legend at 0x21094ab0eb8>



```
In [153]: get_acc_rfc(min_samples_split=35)
```

Out[153]: [0.8491171749598716, 0.8432835820895522]